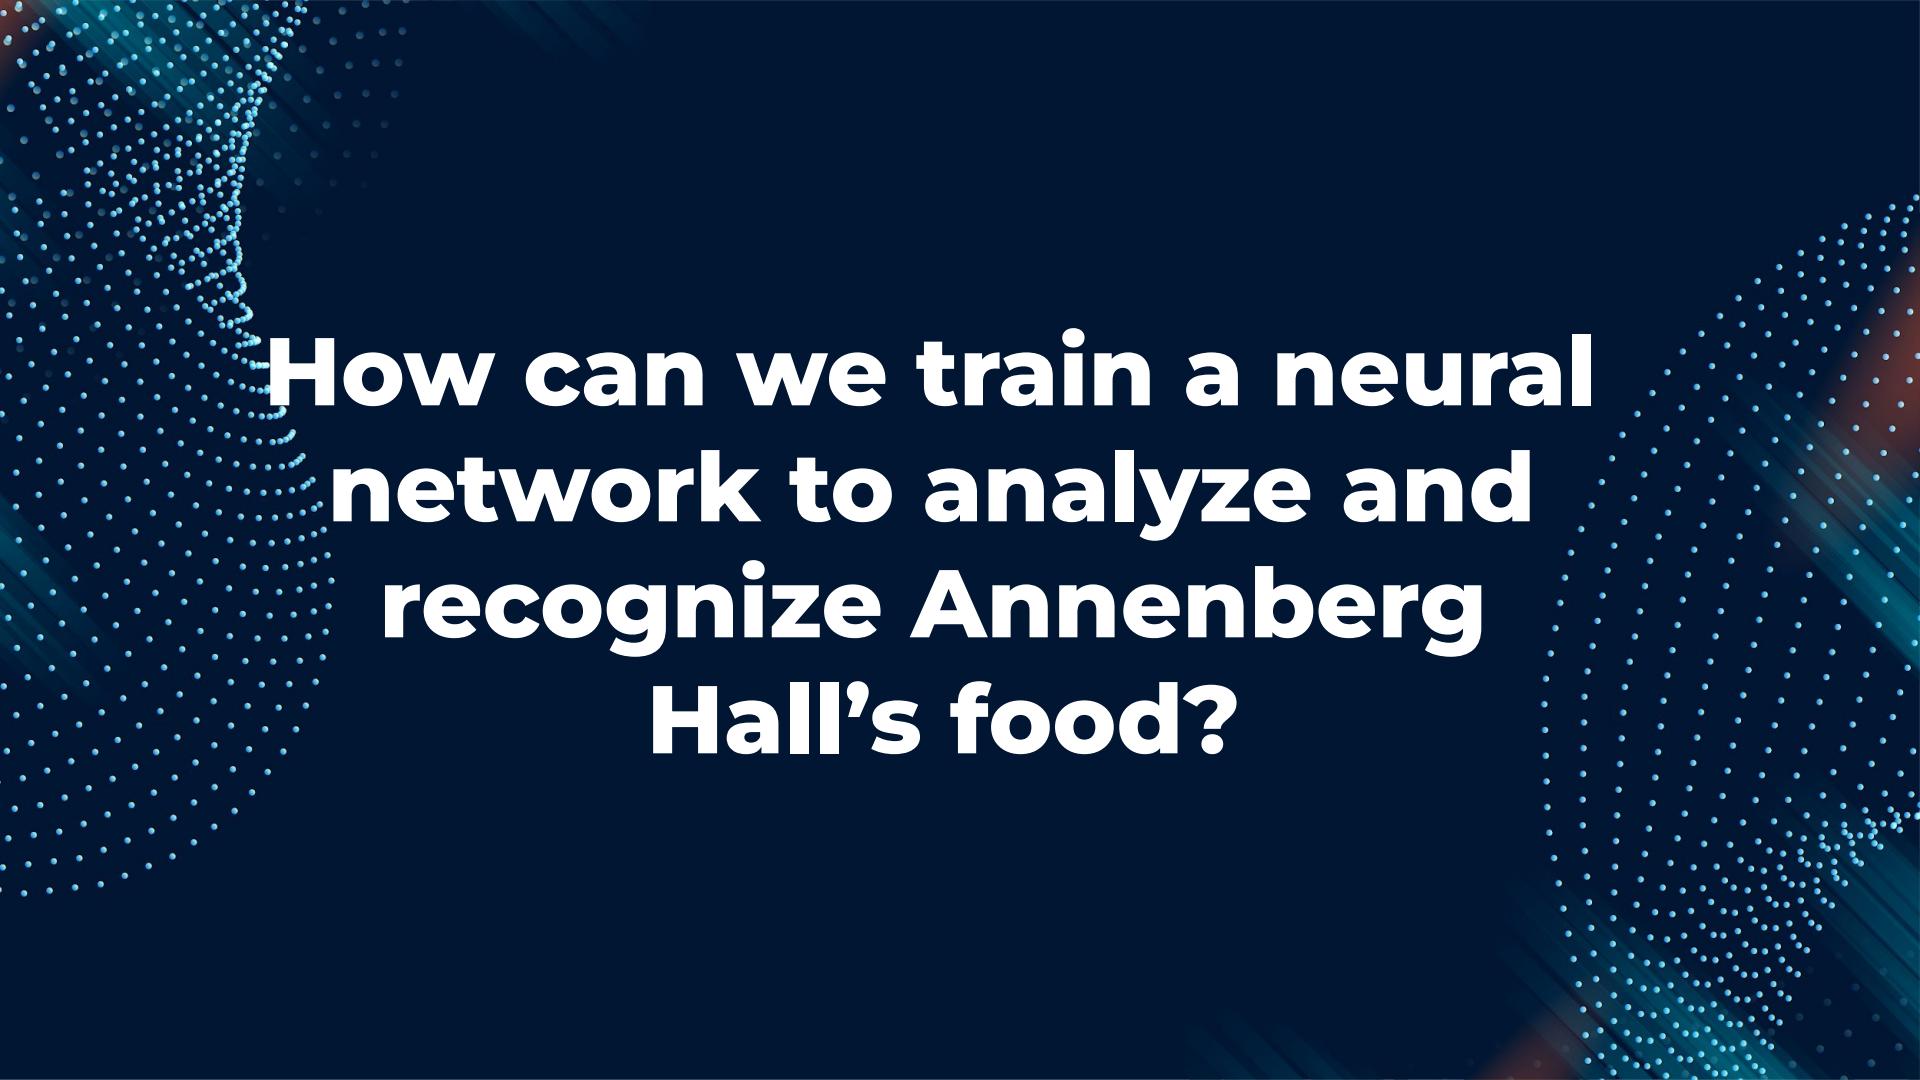




# (Annen) BergNet

**Linda Zeng**

Harvard University  
Pre-College Program  
Presented on August 2, 2024



**How can we train a neural  
network to analyze and  
recognize Annenberg  
Hall's food?**

# Motivation

## This Class

- Ubiquitous part of everyone **here**'s life
- Close to us (literally)
- Fun idea

## In General

- Nutrition recommendation
- Balancing meals
- Monitor food products for food processing
- Machine understanding of images in our world

# Task



Classify

Protein  
**Grains**  
Vegetable  
Fruit  
Soup  
Other

224 x 224

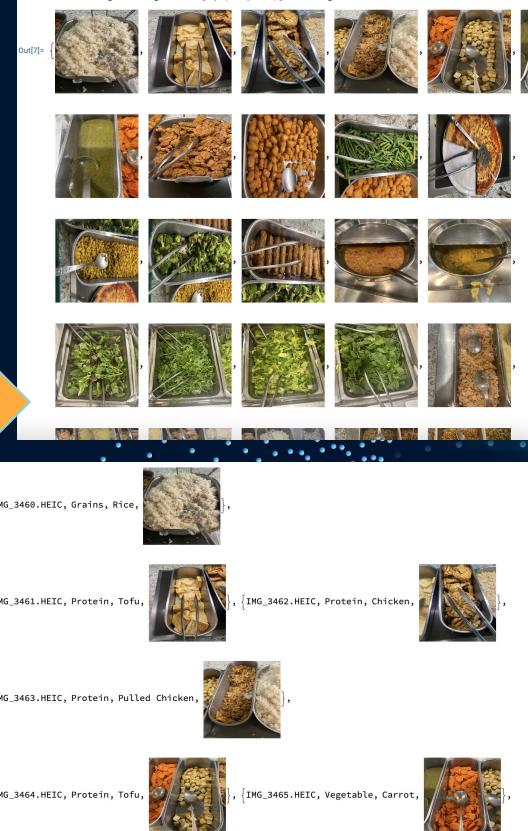
# Data Collection



# Data Labeling

	A	B	
1	Path	Category	Name
2	<a href="#">IMG_3460.HEIC</a>	Grains	Rice
3	<a href="#">IMG_3461.HEIC</a>	Protein	Tofu
4	<a href="#">IMG_3462.HEIC</a>	Protein	Chicken
5	<a href="#">IMG_3463.HEIC</a>	Protein	Pulled Chicken
6	<a href="#">IMG_3464.HEIC</a>	Protein	Tofu
7	<a href="#">IMG_3465.HEIC</a>	Vegetable	Carrot
8	<a href="#">IMG_3466.HEIC</a>	Soup	Green Curry
9	<a href="#">IMG_3467.HEIC</a>	Protein	Fried Chicken
10	<a href="#">IMG_3468.HEIC</a>	Grains	Tater Tots
11	<a href="#">IMG_3469.HEIC</a>	Vegetable	Green Beans
12	<a href="#">IMG_3470.HEIC</a>	Other	Pizza
13	<a href="#">IMG_3471.HEIC</a>	Vegetable	Lima Beans
14	<a href="#">IMG_3472.HEIC</a>	Vegetable	Broccoli
15	<a href="#">IMG_3473.HEIC</a>	Protein	Sausage
16	<a href="#">IMG_3474.HEIC</a>	Soup	Tomato Basil Soup
17	<a href="#">IMG_3475.HEIC</a>	Soup	Chicken Tortilla Soup
18	<a href="#">IMG_6164.HEIC</a>	Vegetable	Spring Mix
19	<a href="#">IMG_6165.HEIC</a>	Vegetable	Baby Arugula
20	<a href="#">IMG_6166.HEIC</a>	Vegetable	Lettuce
21	<a href="#">IMG_6167.HEIC</a>	Vegetable	Spinach
22	<a href="#">IMG_6168.HEIC</a>	Protein	Tuna

# Data Loading



# (Surprising) Challenge

How to add images to dataset? → Transpose!

```
:= newMatrix = Transpose[Join[Transpose[data], {resizedimages}]]
```

```
]:= {{IMG_3460.HEIC, Grains, Rice, },
```



```
{IMG_3461.HEIC, Protein, Tofu, }, {IMG_3462.HEIC, Protein, Chicken, }
```



```
{IMG_3463.HEIC, Protein, Pulled Chicken, },
```



```
{IMG_3464.HEIC, Protein, Tofu, }, {IMG_3465.HEIC, Vegetable, Carrot, }
```



# Data Scarcity Problem

**67** Samples of  $224 \times 224$  pixels.

45 training samples

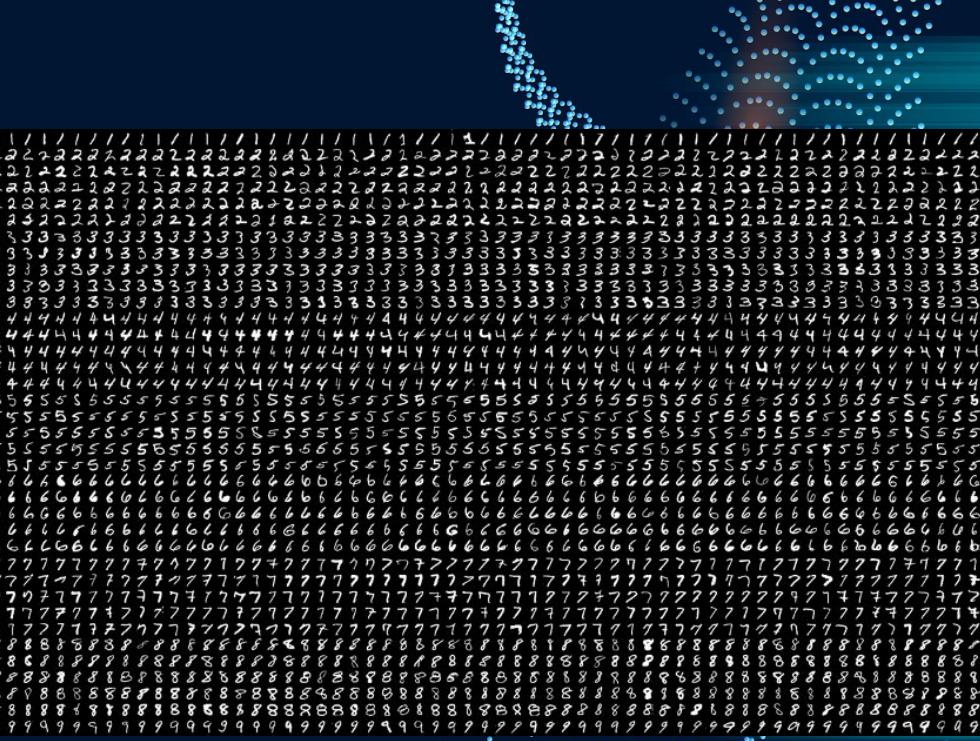
22 test samples

```
In[6]:= Length[images]
```

```
Out[6]= 67
```



VS





VS

# Fake it til you make it!



Flip



Rotate



Translate



Darken



Add noise

Rotate



Brighten

Rotate

# Data Scarcity Problem

**67** Samples of  $224 \times 224$  pixels.

45 training samples

22 test samples

```
In[6]:= Length[images]
```

```
Out[6]= 67
```

# Data Scarcity Problem

**67** Samples of  $224 \times 224$  pixels.

45 training samples

22 test samples



**765 training samples**

```
In[6]:= Length[images]
```

```
Out[6]= 67
```

# Training

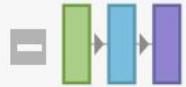
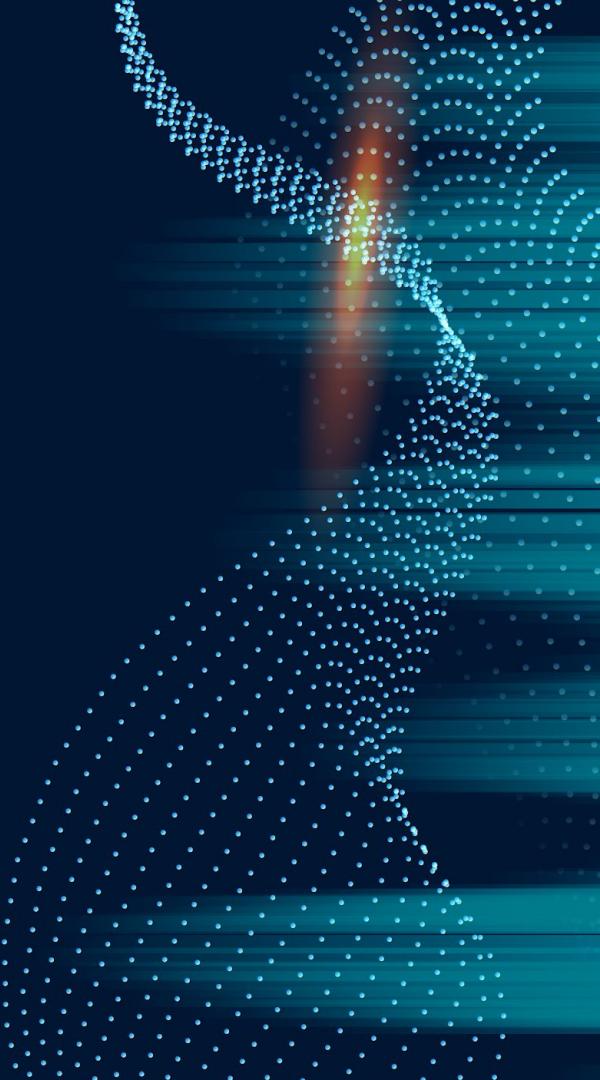


		image
1	Input	array (size: 3×224×224)
2	ConvolutionLayer	array (size: 16×222×222)
3	Ramp	array (size: 16×222×222)
4	PoolingLayer	array (size: 16×111×111)
5	ConvolutionLayer	array (size: 32×109×109)
6	Ramp	array (size: 32×109×109)
7	PoolingLayer	array (size: 32×54×54)
8	FlattenLayer	vector (size: 93312)
9	LinearLayer	vector (size: 128)
10	Ramp	vector (size: 128)
11	LinearLayer	vector (size: 6)
	SoftmaxLayer	vector (size: 6)
	Output	class



# Training



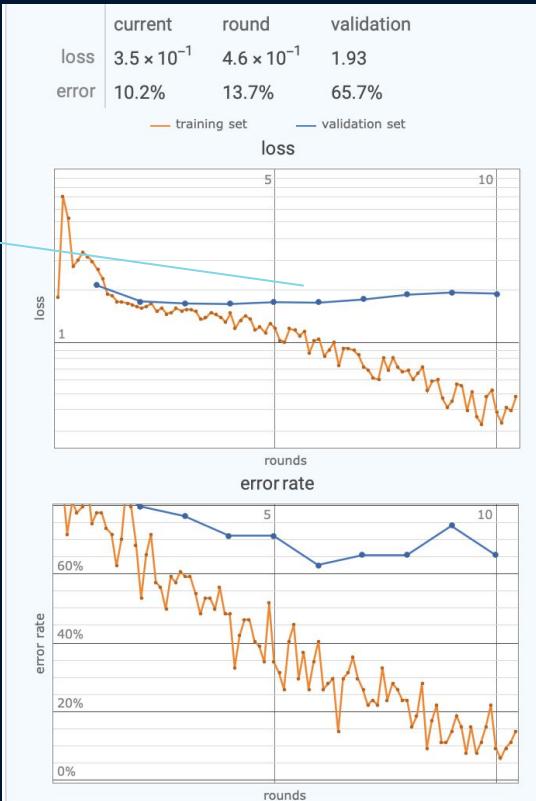
		image
1	Input	array (size: 3×224×224)
2	ConvolutionLayer	array (size: 16×222×222)
3	Ramp	array (size: 16×222×222)
4	PoolingLayer	array (size: 16×111×111)
5	ConvolutionLayer	array (size: 32×109×109)
6	Ramp	array (size: 32×109×109)
7	PoolingLayer	array (size: 32×54×54)
8	FlattenLayer	vector (size: 93312)
9	LinearLayer	vector (size: 128)
10	Ramp	vector (size: 128)
11	LinearLayer	vector (size: 6)
	SoftmaxLayer	vector (size: 6)
	Output	class

98.34% on training set

**48.57% on test set**

# Overfitting

Validation loss starts going up



# Training

No Dropout

98.34% on training set

**48.57%** on test set

Dropout 0.4

62.68% on train

**28.57%** on test

Dropout 0.1

71.3% on train

**37.14%** on test



# Why?

Very difficult task

```
[121]:= trainedModel[
```



```
, {"TopProbabilities", 3}]
```

```
t[121]= {Fruit → 0.00119143, Soup → 0.0524874, Vegetable → 0.945291}
```

```
[122]:= trainedModel[
```



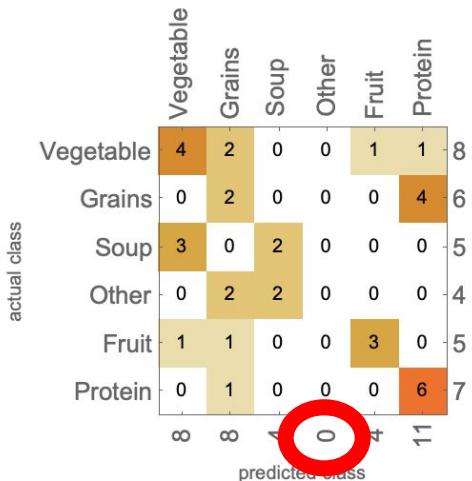
```
, {"TopProbabilities", 3}]
```

```
t[122]= {Other → 0.00118109, Protein → 0.00547433, Vegetable → 0.991891}
```

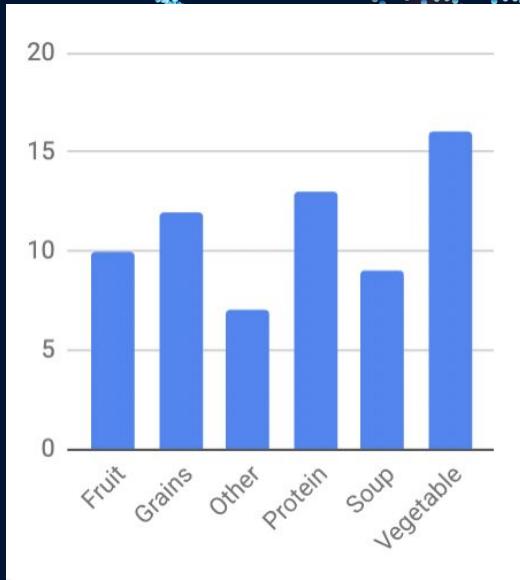
```
= ClassifierMeasurements[trainedBergNet, testData]
```

### Classifier Measurements

Classifier method	Net
Number of test examples	35
Accuracy	(49. ± 9.)%
Accuracy baseline	(23. ± 7.)%
Geometric mean of probabilities	0.118 ± 0.054
Mean cross entropy	2.14 ± 0.44
Single evaluation time	4.98 ms/example
Batch evaluation speed	330. examples/s

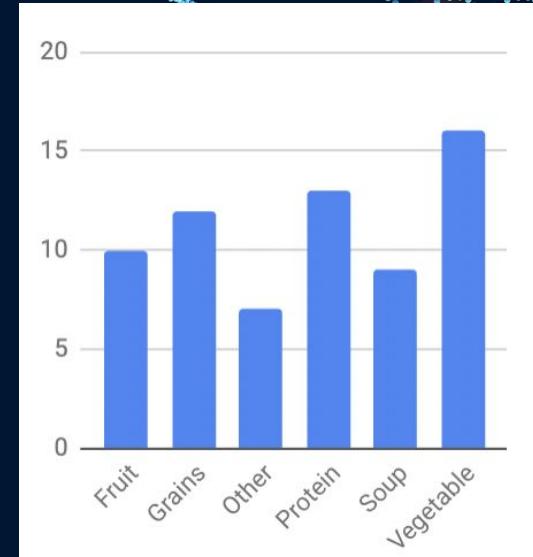


# 1) too little 'Other' samples



## 1) too little 'Other' samples

## 2) 'Other' samples are too random



1,



→ 1,



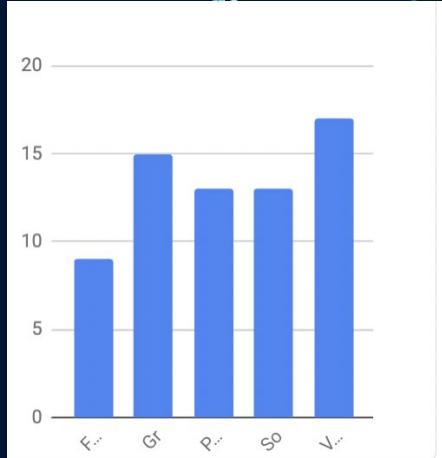
→ 1,



→ 1,

# Removing “Other”

Going back to the **beginning**



A dot plot visualization on the right side of the slide, showing a dense cluster of dots forming a shape that suggests a spiral or a path through the data points.

VALUE	FREQUENCY
Vegetable	17
Grains	15
Soup	13
Protein	13
Fruit	9

# BergNet Results

98.34%

99.34% on training set

**48.57%**

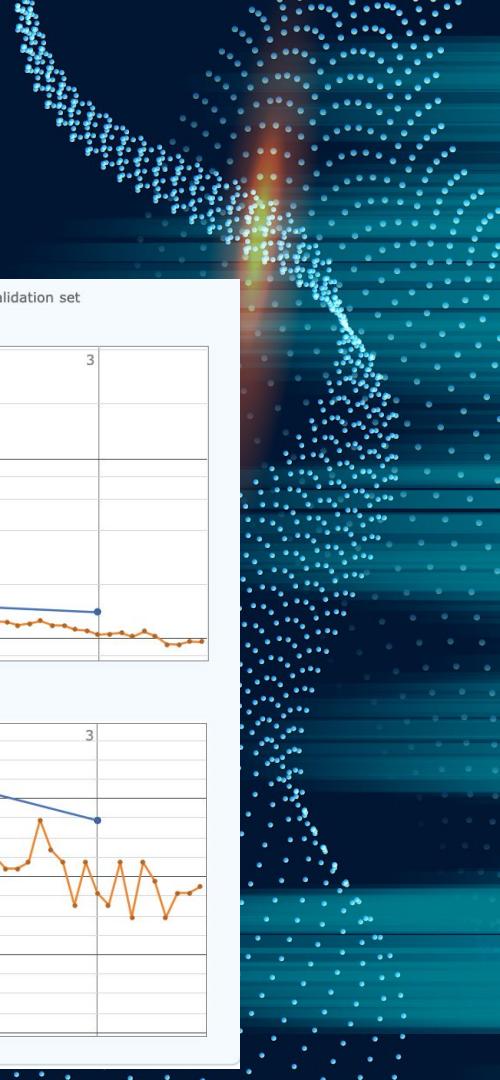
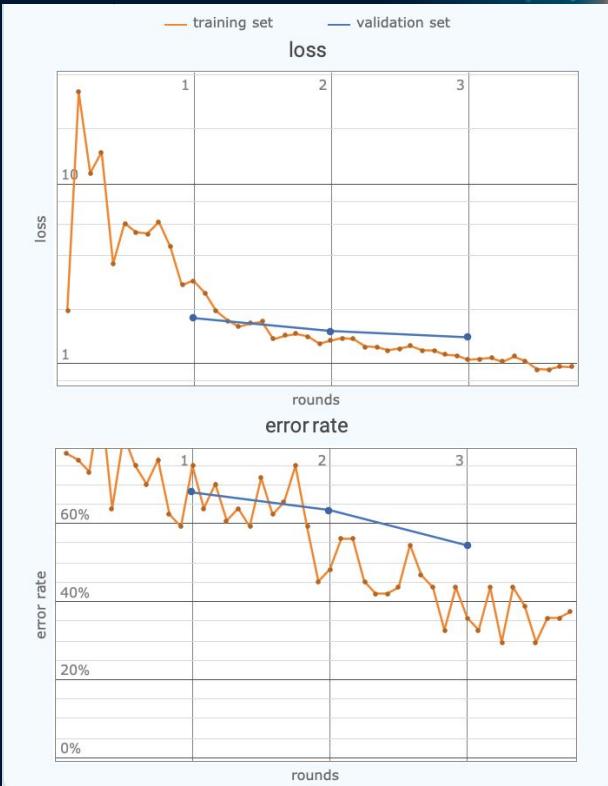
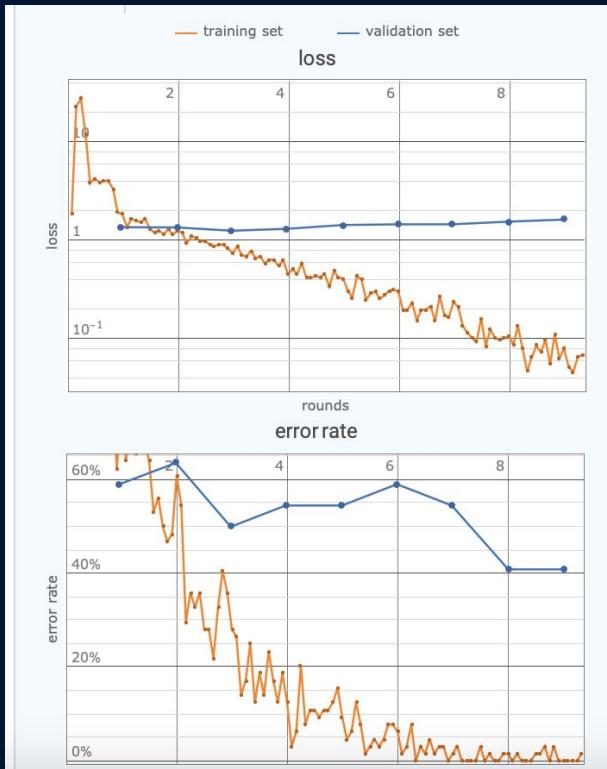
**59% on test set**

Removed 'Other' class

	Input	image
1	ConvolutionLayer	array (size: 3×224×224)
2	Ramp	array (size: 16×222×222)
3	PoolingLayer	array (size: 16×111×111)
4	ConvolutionLayer	array (size: 32×109×109)
5	Ramp	array (size: 32×109×109)
6	PoolingLayer	array (size: 32×54×54)
7	FlattenLayer	vector (size: 93312)
8	LinearLayer	vector (size: 128)
9	Ramp	vector (size: 128)
10	LinearLayer	vector (size: 5)
11	SoftmaxLayer	vector (size: 5)
	Output	class

# Training (trial 2)

Added dropout 0.1



# BergNet Results

98.34%

99.34% on training set

**48.57%**

**59% on test set**

Removed 'Other' class

91.50% on training set

**63.63% on test set**

Dropout 0.1

# Failures

<b>Best</b>	<b>91.50%</b>	<b>63.63%</b>
Dropout 0.2	84.84%	59.09%
Add early stopping	77.52%	50%
Add L2 Regularization + SGD	59%	40.9%
Add L2 Regularization + Adam	93.07%	54.54%
Add Batch Norm	27%	27%

## Overall results

	# Data points	Accuracy on Wednesday Test Data
No Training	0	23%
BergNet (no data augmentation)	45	50%
BergNet	765	<b>64%</b>

## Examples (BergNet)



Predicted as Grains



Predicted as Protein



Predicted as Grains



Predicted as Soup

# Examples



Predicted as Grains



Predicted as Protein



Predicted as Grains



Predicted as Soup



# **What else can we do?**

- Constrained by limited data



# What else can we do?

- Constrained by data



## Rely on others

- There are already models trained on **similar** tasks



# Fine-tune EfficientNet

Trained to identify images

- trained on 1.2 million images
- Already optimized parameters

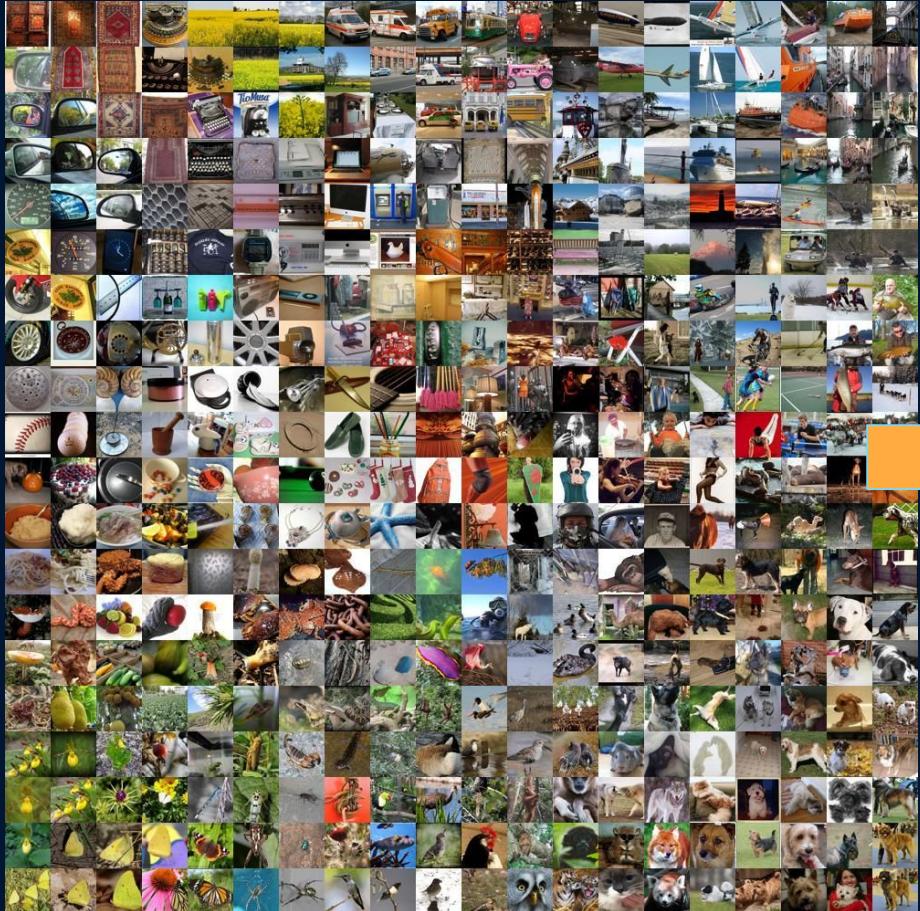
We can modify it for our tasks

- train on our data

Replace with 5 classes →



NetChain[ ]		
	Input	image
stem_conv	ConvolutionLayer	array (size: 3x224x224)
stem_bn	BatchNormalizationLayer	array (size: 32x112x112)
stem_activation	LogisticSigmoid[x]x	array (size: 32x112x112)
block1a	NetChain (6 nodes)	array (size: 16x112x112)
block2a	NetChain (9 nodes)	array (size: 24x56x56)
block2b	NetGraph (11 nodes)	array (size: 24x56x56)
block3a	NetChain (9 nodes)	array (size: 40x28x28)
block3b	NetGraph (11 nodes)	array (size: 40x28x28)
block4a	NetChain (9 nodes)	array (size: 80x14x14)
block4b	NetGraph (11 nodes)	array (size: 80x14x14)
block4c	NetGraph (11 nodes)	array (size: 80x14x14)
block5a	NetChain (9 nodes)	array (size: 112x14x14)
block5b	NetGraph (11 nodes)	array (size: 112x14x14)
block5c	NetGraph (11 nodes)	array (size: 112x14x14)
block6a	NetChain (9 nodes)	array (size: 192x7x7)
block6b	NetGraph (11 nodes)	array (size: 192x7x7)
block6c	NetGraph (11 nodes)	array (size: 192x7x7)
block6d	NetGraph (11 nodes)	array (size: 192x7x7)
block7a	NetChain (9 nodes)	array (size: 320x7x7)
top_conv	ConvolutionLayer	array (size: 1280x7x7)
top_bn	BatchNormalizationLayer	array (size: 1280x7x7)
top_activation	LogisticSigmoid[x]x	array (size: 1280x7x7)
avg_pool	AggregationLayer	vector (size: 1280)
top_dropout	DropoutLayer	vector (size: 1280)
<b>predictions</b>	<b>LinearLayer</b>	<b>vector (size: 1000)</b>
<b>predictions_activation</b>	<b>SoftmaxLayer</b>	<b>vector (size: 1000)</b>
	<b>Output</b>	<b>class</b>



Some but not all are food pics

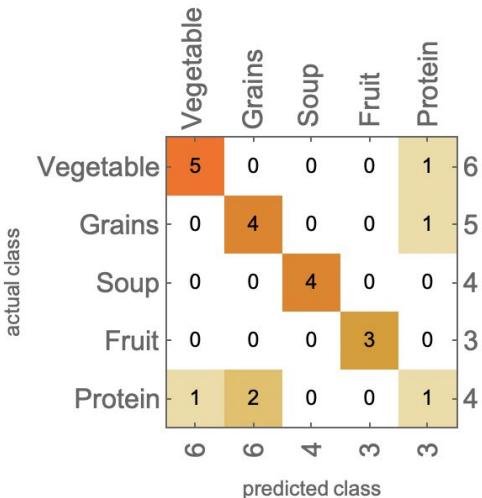


Add on these

# Gets 77% accuracy!

## Classifier Measurements

Classifier method	Net
Number of test examples	22
Accuracy	(77. ± 9.)%
Accuracy baseline	(27. ± 10.)%
Geometric mean of probabilities	$0.326 \pm 0.13$
Mean cross entropy	$1.12 \pm 0.39$
Single evaluation time	59. ms/example
Batch evaluation speed	15.1 examples/s



## Overall results

	# Data points	Accuracy on Wednesday Test Data
No Training	0	23%
BergNet (no data augmentation)	45	50%
BergNet	765	<b>64%</b>

## Overall results

	# Data points	Accuracy on Wednesday Test Data
No Training	0	23%
BergNet (no data augmentation)	45	50%
BergNet	765	64%
EfficientNet + ImageNet	0	14%
EfficientNet + ImageNet (Finetuned)	765	<b>77%</b>

# Examples (Finetuned EfficientNet)



Predicted as ~~Grains~~

Predicted as **Fruit**



Predicted as Protein

Predicted as **Protein**



Predicted as ~~Grains~~

Predicted as **Protein**



Predicted as ~~Soup~~

Predicted as **Vegetable**



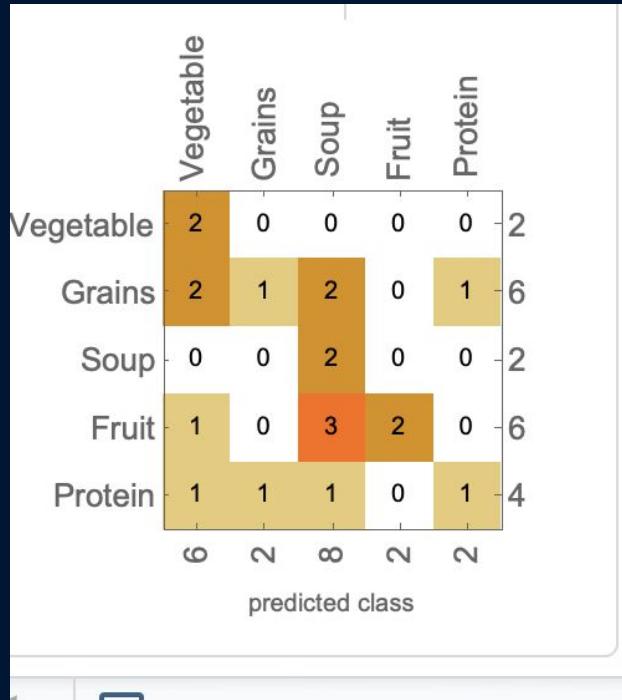
**Time to test out on  
Thursday now!**

## Overall results (with Thursday)

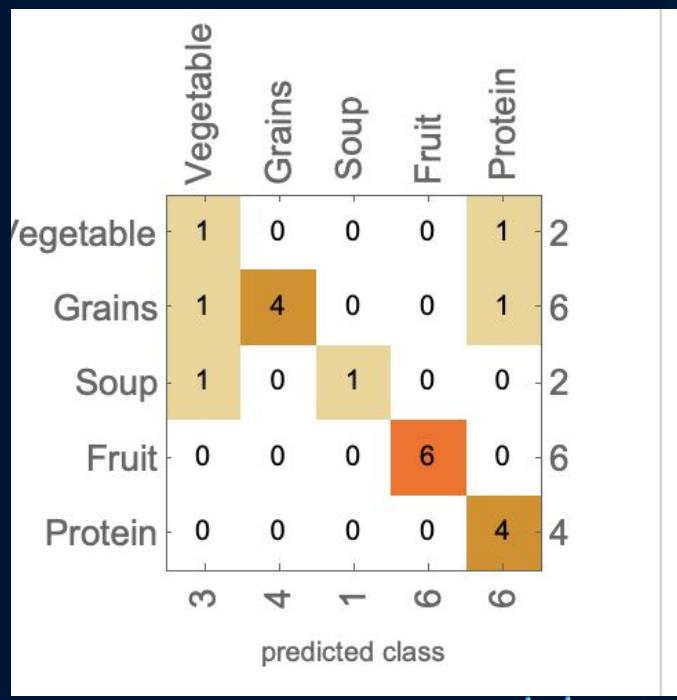
	# Data points	Accuracy on Wednesday Test Data	Accuracy on Thursday Test
No Training	0	23%	25%
BergNet (no data augmentation)	45	50%	30%
BergNet	765	64%	40%
EfficientNet + ImageNet	0	14%	35%
EfficientNet + ImageNet (finetuned)	765	<b>77%</b>	<b>80%</b>

# Thursday

BertNet

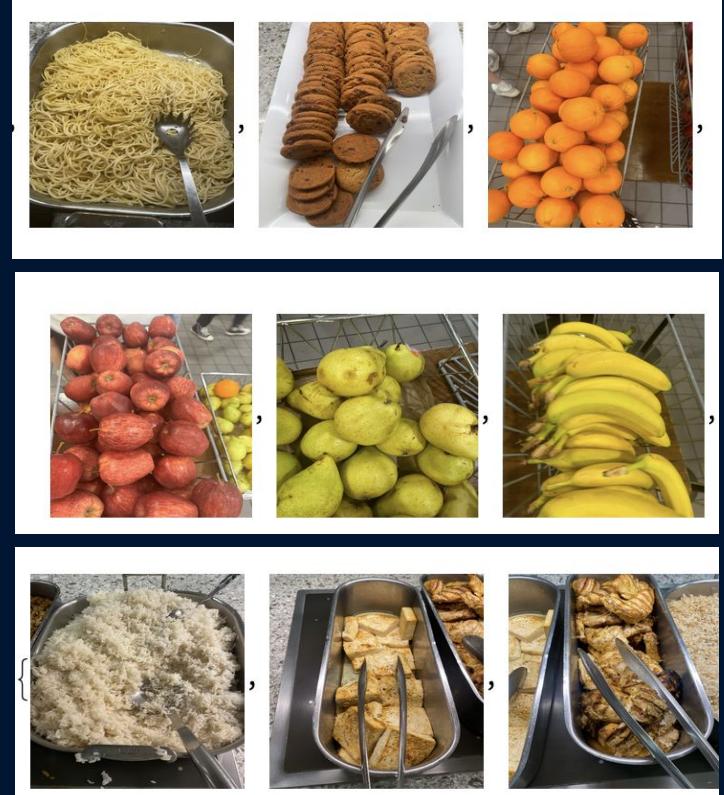
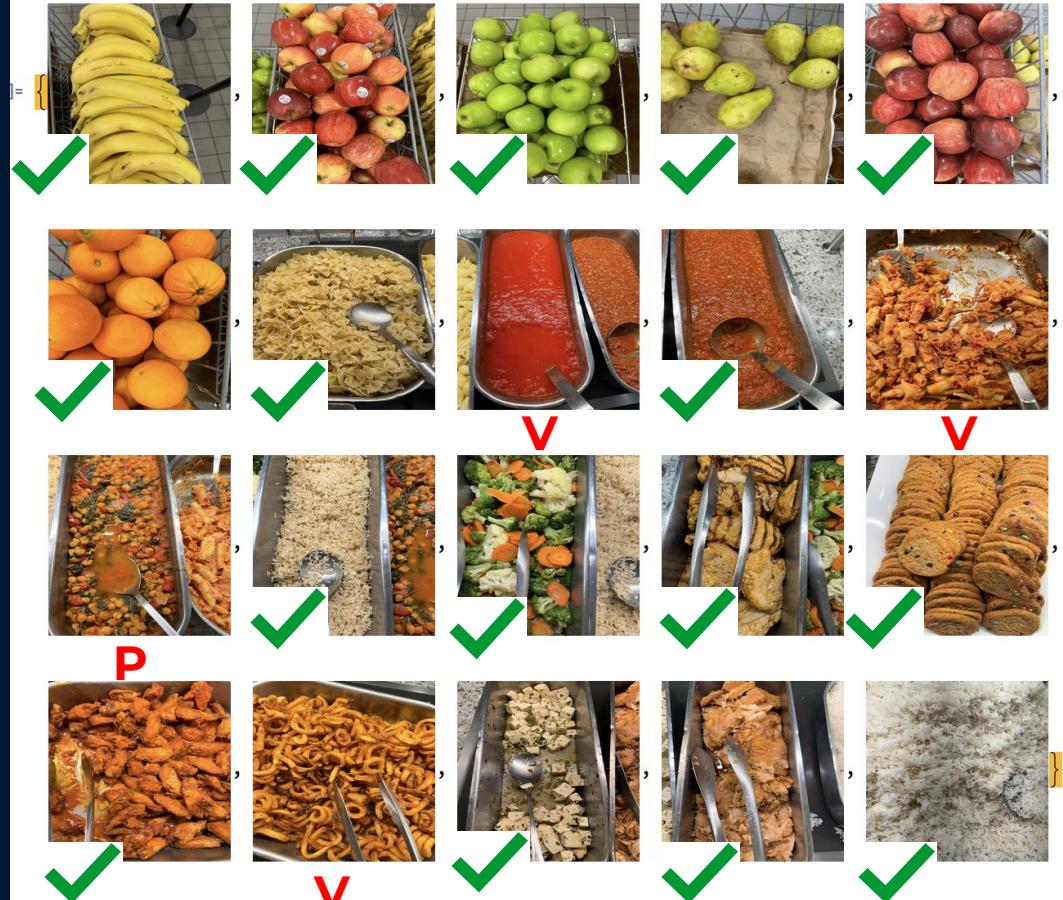


Finetuned EfficientNet

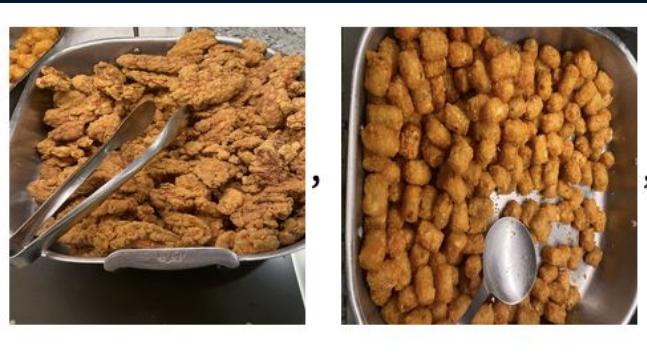




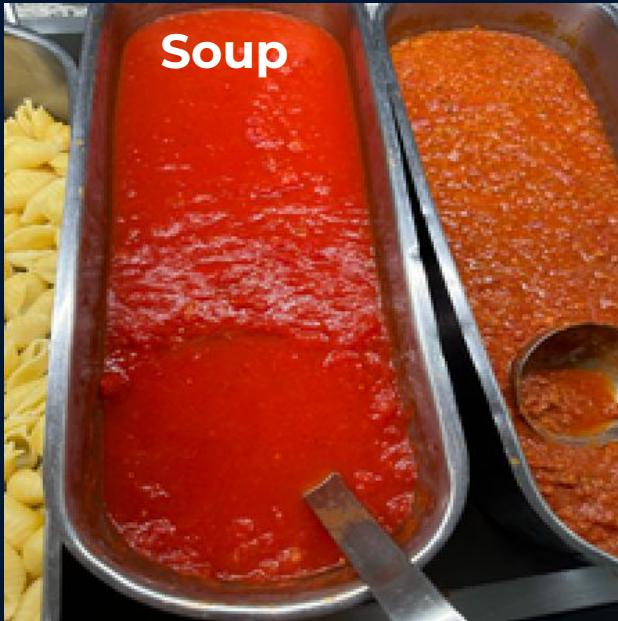
# Everyday foods



## Exceptions



# Exceptions

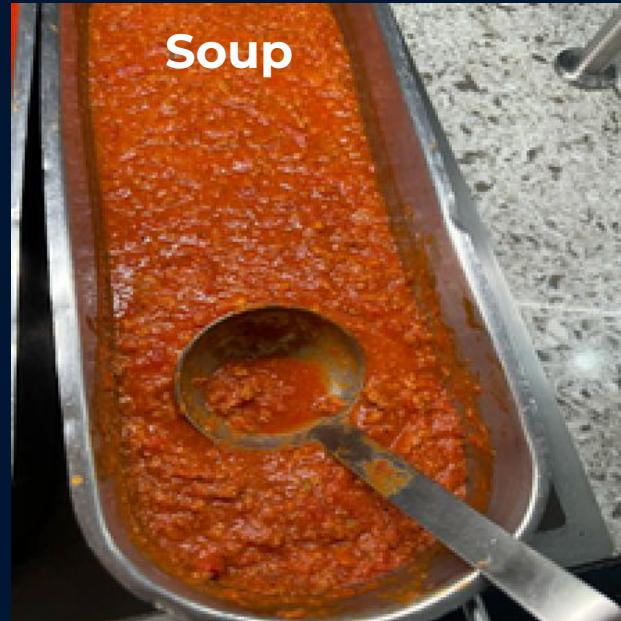


Grains -> 0.106108

Soup -> 0.375381

**Vegetable ->**

**0.47878**



Grains -> 0.349892

**Soup -> 0.412876**

Vegetable -> 0.148721

# Conclusion

- We successfully got 77% on Wednesday and 80% on Thursday
- There are ways to work with low amounts of data
  - **Fake it til you make it:** Data augmentation
  - **Rely on others:** Fine tuning
- The more consistent food Annenberg makes, the better BergNet performs